

practical5-piyusha

April 3, 2025

0.1 Piyusha Supe 23CO315

Practical 5: Data Analytics II 1. Implement logistic regression using Python/R to perform classification on Social_Network_Ads.csv dataset. 2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score,
                           classification_report, precision_score, recall_score, f1_score
```

```
[ ]: from google.colab import files
files.upload()
```

```
[3]: df = pd.read_csv("Social_Network_Ads.csv")
```

```
[4]: # check the dataset for missing values and perform EDA
df.head()
```

```
[4]:   User ID  Gender  Age  EstimatedSalary  Purchased
0    15624510    Male   19        19000          0
1    15810944    Male   35        20000          0
2    15668575  Female   26        43000          0
3    15603246  Female   27        57000          0
4    15804002    Male   19        76000          0
```

```
[5]: df.tail()
```

```
[5]:   User ID  Gender  Age  EstimatedSalary  Purchased
395  15691863  Female   46        41000          1
396  15706071    Male   51        23000          1
397  15654296  Female   50        20000          1
398  15755018    Male   36        33000          0
```

```
399 15594041 Female 49          36000      1
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   User ID           400 non-null    int64  
 1   Gender             400 non-null    object  
 2   Age                400 non-null    int64  
 3   EstimatedSalary   400 non-null    int64  
 4   Purchased          400 non-null    int64  
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
[7]: df.describe(include="all")
```

```
[7]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
count	4.000000e+02	400	400.000000	400.000000	400.000000
unique	NaN	2	NaN	NaN	NaN
top	NaN	Female	NaN	NaN	NaN
freq	NaN	204	NaN	NaN	NaN
mean	1.569154e+07	NaN	37.655000	69742.500000	0.357500
std	7.165832e+04	NaN	10.482877	34096.960282	0.479864
min	1.556669e+07	NaN	18.000000	15000.000000	0.000000
25%	1.562676e+07	NaN	29.750000	43000.000000	0.000000
50%	1.569434e+07	NaN	37.000000	70000.000000	0.000000
75%	1.575036e+07	NaN	46.000000	88000.000000	1.000000
max	1.581524e+07	NaN	60.000000	150000.000000	1.000000

```
[8]: df.shape
```

```
[8]: (400, 5)
```

```
[9]: df.size
```

```
[9]: 2000
```

```
[10]: df.ndim
```

```
[10]: 2
```

```
[11]: df.columns
```

```
[11]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'],  
          dtype='object')
```

```
[12]: df.isnull()
```

```
[12]:      User ID  Gender    Age  EstimatedSalary  Purchased  
0        False   False  False           False  False  
1        False   False  False           False  False  
2        False   False  False           False  False  
3        False   False  False           False  False  
4        False   False  False           False  False  
..        ...     ...  ...           ...  ...  
395       False   False  False           False  False  
396       False   False  False           False  False  
397       False   False  False           False  False  
398       False   False  False           False  False  
399       False   False  False           False  False
```

[400 rows x 5 columns]

```
[13]: df.isna().sum()
```

```
[13]: User ID      0  
Gender       0  
Age          0  
EstimatedSalary  0  
Purchased     0  
dtype: int64
```

```
[15]: #select the features and target variable  
# Selecting independent variables (features) and dependent variable (target)  
X = df[['Age', 'EstimatedSalary']] # Assuming these are features  
y = df['Purchased'] # Target variable
```

```
[16]: # Splitting dataset into training (80%) and testing (20%) sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
                                                 random_state=42)  
  
print("\nTraining Set Size:", X_train.shape)  
print("Testing Set Size:", X_test.shape)
```

Training Set Size: (320, 2)
Testing Set Size: (80, 2)

```
[17]: # Train the Logistic Regression model  
model = LogisticRegression()
```

```
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)
```

```
[18]: accuracy = accuracy_score(y_test, y_pred)
print(f"\nAccuracy: {accuracy:.4f}")
```

Accuracy: 0.8875

```
[19]: # Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Extract values
TN, FP, FN, TP = cm.ravel()

# Print confusion matrix and extracted values
print("\nConfusion Matrix:")
print(cm)
print(f"\nTrue Positives (TP): {TP}")
print(f"False Positives (FP): {FP}")
print(f"True Negatives (TN): {TN}")
print(f"False Negatives (FN): {FN}")
```

Confusion Matrix:

```
[[50  2]
 [ 7 21]]
```

True Positives (TP): 21
False Positives (FP): 2
True Negatives (TN): 50
False Negatives (FN): 7

```
[21]: precision = precision_score(y_test, y_pred)
print(f"\nPrecision: {precision:.4f}")
```

Precision: 0.9130

```
[22]: recall = recall_score(y_test, y_pred)
print(f"\nRecall: {recall:.4f}")
```

Recall: 0.7500

```
[23]: error_rate = 1 - accuracy
print(f"\nError Rate: {error_rate:.4f}")
```

Error Rate: 0.1125


```
[24]: f1 = f1_score(y_test, y_pred)
print(f"\nF1 Score: {f1:.4f}")
```

F1 Score: 0.8235

```
[25]: # Display detailed classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Classification Report:				
	precision	recall	f1-score	support
0	0.88	0.96	0.92	52
1	0.91	0.75	0.82	28
accuracy			0.89	80
macro avg	0.90	0.86	0.87	80
weighted avg	0.89	0.89	0.88	80

```
[26]: # Visualizing confusion matrix
plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Not Purchased", "Purchased"], yticklabels=["Not Purchased", "Purchased"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()
```



```
[31]: import matplotlib.pyplot as plt
import seaborn as sns

# Scatter plot of test set predictions
# Accessing DataFrame columns using column names
plt.scatter(X_test['Age'], X_test['EstimatedSalary'], c=y_pred,
            cmap='coolwarm', edgecolors='k')

# Labels and title
plt.xlabel('Age (Standardized)')
plt.ylabel('Estimated Salary (Standardized)')
plt.title('Logistic Regression Classification')

plt.show()
```

Logistic Regression Classification

